

StationSim code acceleration

Project RADDISH

Eleftherios Avramidis
Research Software Engineering
Research Computing Services
University of Cambridge

StationSim

- Simulates crowds in real time with agent-based modelling and a particle filter (Nicolas Malleson, et al. 2019)
- Written in Python
- Good code quality
- No tests
- Jupyter notebooks examples
- Used with typical size of 40 agents and 10,000 particles

Project aim and organisation

- Multi-core/node and GPU implementation of the StationSim model and particle filter
- Stage 1 (done):
 - Familiarisation with model and Python code
 - Establishment of implementation strategy
- Stage 2 (in progress): Implementation using C++ with OpenMP and MPI
- Stage 3: GPU implementation

Extra aims

- Maintainable code
- Addition of tests and CI
- Write instructions for
 - Building
 - Running
 - Testing

StationSimCpp

- Private GitHub repository
- C++ shared library
- C++17 with OpenMP and MPI
- CMake to build, test and package the library
- CI using GitLab
- Needs more test (unit, integration)
- HDF5 for file output

Early runtime/speedup results – one CPU core

StationSim single model run with default parameters

Population	Python (s)	C++ (s)	Speedup (Python/C++)
100	4.58	0.034	134.7059
200	13.03	0.14	93.07143
400	40.51	0.74	54.74324
600	86.19	2.1	41.04286
800	153.54	4.55	33.74505
1000	265.71	7.67	34.64276

Preliminary results

Early runtime results – 6 CPU cores using OpenMP

**StationSim multiple model runs with default parameters
and population size of 400**

Model instances (i.e. particles)	Runtime
100	11.1435
200	22.1494
400	40.3555
600	61.8257

Preliminary results

Next steps

- Improve MPI multi-node implementation
- Implement simple particle filter
- Implement HDF5 outputs code for particle filter
- Performance profiling
- GPU code
- Improve repository and documentation
- Finalise code API